

## Milestone 2

Jasper Evans, Krehl Kasayan, Paul Kiefer, Pablo Suarez

INST737

4/8/2024

**Introduction:** Before we could begin building and assessing our predictive models, we had to transform raw American Community Survey data into scaled metrics that can be compared between census tracts. For context, the American Community Survey is an annual survey of a sample of US residents at various levels of geography. Because it takes time to collect those samples, survey results are provided in five-year spans.

Those variables are:

1. avg\_bed: The average number of beds in a housing unit at the tract level as of the 2016-2020 ACS.
2. tract\_homevalue\_2020: The mean self-reported home value at the tract level as of the 2016-2020 ACS.
3. tract\_medage\_2020: The median age of residents at the tract level as of the 2016-2020 ACS.
4. tract\_medincome\_2020: The median income of households at the tract level as of the 2016-2020 ACS.
5. tract\_medincome\_2010: The median income of households at the tract level as of the 2006-2010 ACS.
6. foreclosure\_pc\_2010: The number of foreclosures per 1,000 residents at the tract level during the height of the Great Recession from 2008-2010.
7. pct\_built\_2020\_later: The percentage of housing units in a tract built in or after 2020.
8. pct\_built\_2010\_2019: The percentage of housing units in a tract built between 2010-2019.
9. pct\_built\_2000\_2009: The percentage of housing units in a tract built between 2000-2009.
10. pct\_built\_1990\_1999: The percentage of housing units in a tract built between 1990-1999.
11. pct\_built\_1980\_1989: The percentage of housing units in a tract built between 1980-1989.
12. pct\_built\_1970\_1979: The percentage of housing units in a tract built between 1970-1979.
13. pct\_built\_pre\_1960: The percentage of housing units in a tract built prior to 1960.
14. pct\_0\_bd: The percentage of housing units with no separate bedroom (i.e. studios) in a tract.
15. pct\_1\_bd: The percentage of housing units with one bedroom in a tract.
16. pct\_2\_bd: The percentage of housing units with two bedrooms in a tract.
17. pct\_3\_bd: The percentage of housing units with 3 bedrooms in a tract.
18. pct\_4\_more\_bd: The percentage of housing units with 4 or more bedrooms in a tract.
19. poverty\_2010: The percentage of a tract's population with incomes below the federal poverty line as of the 2006-2010 ACS.

20. poverty\_2020: The percentage of a tract's population with incomes below the federal poverty line as of the 2016-2020 ACS.
21. nhwhite\_2010: The percentage of a tract's population that self-identified as non-Hispanic white in the 2006-2010 ACS.
22. nhwhite\_2020: The percentage of a tract's population that self-identified as non-Hispanic white in the 2016-2020 ACS.
23. mortgaged\_2010: The percentage of all housing units in a tract with a mortgage or similar loan in the 2006-2010 ACS.
24. mortgaged\_2015: The percentage of all housing units in a tract with a mortgage or similar loan in the 2011-2015 ACS.
25. mortgaged\_2020: The percentage of all housing units in a tract with a mortgage or similar loan in the 2016-2020 ACS.
26. ownoccupied\_2010: The percentage of all housing units in a tract that were owner-occupied in the 2006-2010 ACS.
27. ownoccupied\_2015: The percentage of all housing units in a tract that were owner-occupied in the 2011-2015 ACS.
28. ownoccupied\_2020: The percentage of all housing units in a tract that were owner-occupied in the 2016-2020 ACS.
29. mortgage\_change\_2010\_2015: The change in the percentage of all housing units in a tract with a mortgage or similar loan between the 2006-2010 ACS and 2011-2015 ACS.
30. mortgage\_change\_2015\_2020: The change in the percentage of all housing units in a tract with a mortgage or similar loan between the 2011-2015 ACS and 2016-2020 ACS.
31. mortgage\_change\_2010\_2020: The change in the percentage of all housing units in a tract with a mortgage or similar loan between the 2006-2010 ACS and 2016-2020 ACS.
32. ownoccupied\_change\_2010\_2015: The change in the percentage of all housing units in a tract that were owner-occupied between the 2006-2010 ACS and 2011-2015 ACS.
33. ownoccupied\_change\_2015\_2020: The change in the percentage of all housing units in a tract that were owner-occupied between the 2011-2015 ACS and 2016-2020 ACS.
34. ownoccupied\_change\_2010\_2020: The change in the percentage of all housing units in a tract that were owner-occupied between the 2006-2010 ACS and 2016-2020 ACS.
35. poverty\_change\_2010\_2020: The change in the percentage of all residents in a tract living below the federal poverty line between the 2006-2010 ACS and 2016-2020 ACS.
36. nhwhite\_change\_2010\_2020: The change in the percentage tract residents who self-identified as non-Hispanic white between the 2006-2010 ACS and 2016-2020 ACS.
37. medincome\_change\_2010\_2015: The change in the reported tract-level median household income between the 2006-2010 ACS and 2011-2015 ACS.
38. medincome\_change\_2015\_2020: The change in the reported tract-level median household income between the 2011-2015 ACS and 2016-2020 ACS.
39. medincome\_change\_2010\_2020: The change in the reported tract-level median household income between the 2006-2010 ACS and 2016-2020 ACS.
40. pop\_change\_pct: The change in the reported tract-level population between the 2006-2010 ACS and 2016-2020 ACS.

These data points are all compared to our response variable (or a version of it):  
foreclosure\_pc\_2020.

That variable represents the number of reported foreclosures in a census tract between 2011-2023 divided by the estimated population of the tract in the 2016-2020 ACS. That may seem like an imperfect point of comparison, and in some ways, it is: ideally, we would work with foreclosures per capita in the most recent year for which we have both foreclosure and population data available.

Unfortunately, the Prince George's County dataset that records foreclosures by address contains very few records dated after 2020 – which may mean some foreclosures have been omitted – so we would have hardly any samples if we were to limit ourselves to recent years.

Instead, we consider all foreclosures between the end of the Great Recession – defined in our case as 2011, though that is up for debate – and the end of our dataset in 2011. This gives us a larger sample of tracts to assess, and we include variables measuring characteristics of each tract at various points in that decade to ensure we are considering changes in a tract's risks as part of our analysis.

We may have too many variables for a sample size of only roughly 170 tracts, and in the next step of this process, we may cut back the number of variables while attempting to increase our sample size.

None of the models attempted in this iteration of the project were ideal, and there are certainly other variables that could be more predictive than what we have available. Federal Home Mortgage Disclosure Act (HMDA) data, for instance, could give us tract-level statistics on the percentage of mortgage loans that were denied over a given time period, the percentage of approved mortgage loans that were categorized as refinancing over a given period of time, and the percentage of mortgage loans given for non-owner-occupied properties over a given time period. Unfortunately, the census tract ID numbers included in the HMDA database are incomplete, and we have not yet found a way to correct them.

Another possible tweak would involve shifting to census block-level analysis, which would increase the size of our sample but decrease the number of statistics available on the residents and housing units in a given block. Block-level statistics are only collected during decennial censuses, and decennial censuses are less comprehensive than the American Community Surveys.

### **Question 1: Linear Regressions**

***Divide your dataset into training and testing sets as we have seen in class and report:***

### ***Linear Regression Parameters:***

Because our dataset included some columns that do not contain independent variables, we began this process by defining the independent variables of interest for our analysis using the following code:

```
variables_of_interest <- c("avg_bed", "tract_homevalue_2020",  
  "tract_medage_2020", "tract_medincome_2020", "tract_medincome_2010",  
  "foreclosure_pc_2010", "pct_built_2020_later", "pct_built_2010_2019",  
  "pct_built_2000_2009", "pct_built_1990_1999", "pct_built_1980_1989",  
  "pct_built_1970_1979", "pct_built_pre_1960", "pct_0_bd", "pct_1_bd",  
  "pct_2_bd", "pct_3_bd", "pct_4_more_bd", "poverty_2010", "poverty_2020",  
  "nhwhite_2010", "nhwhite_2020", "mortgaged_2010", "mortgaged_2015",  
  "mortgaged_2020", "ownoccupied_2010", "ownoccupied_2015",  
  "ownoccupied_2020", "mortgage_change_2010_2015",  
  "mortgage_change_2015_2020", "mortgage_change_2010_2020", "ownoccupied_change  
_2010_2015", "ownoccupied_change_2015_2020", "ownoccupied_change_2010_2020",  
  "poverty_change_2010_2020", "nhwhite_change_2010_2020", "medincome_change_201  
0_2015", "medincome_change_2015_2020", "medincome_change_2010_2020",  
  "pop_change_pct")
```

We then divide our dataset into a training set – a random sample of 70% of our dataset – and a test set made up of a random sample of 30% of our dataset.

***For each independent variable in your model, compute a linear regression with respect to the dependent feature and report:***

We then use the following code to loop through each of the 40 independent variables and calculate the **intercept, correlation coefficient and mean squared error of each independent variable relative to the response variable, as well as whether the independent variable is predictive.**

```
for (variable in variables_of_interest) {  
  # Perform linear regression  
  model <- lm(foreclosure_pc_2020 ~ ., data = pg_foreclosures_train[,  
c(variable, "foreclosure_pc_2020")])  
  
  # Extract required information  
  intercept <- coef(model)[1]  
  coefficient <- coef(model)[2]  
  
  # Check if the coefficient is statistically significant
```

```

p_value <-
summary(model)$coefficients[which(rownames(summary(model)$coefficients) ==
variable), "Pr(>|t|)"]
is_predictive <- ifelse(p_value < 0.05, "Yes", "No")

# Compute residuals
residuals <- resid(model)

# Compute correlation between predicted and real values
correlation <- cor(predict(model),
pg_foreclosures_train$foreclosure_pc_2020 )

# Compute mean square error
mse <- mean((predict(model) - pg_foreclosures_train$foreclosure_pc_2020
)^2)

# Print results
cat("Variable:", variable, "\n")
cat("Intercept:", intercept, "\n")
cat("Coefficient:", coefficient, "\n")
cat("Is it a predictive feature?:", is_predictive, "\n")
cat("Correlation:", correlation, "\n")
cat("Mean Square Error:", mse, "\n")
cat("\n")
}

```

That produces the following results:

```

Variable: avg_bed
Intercept: 19.06023
Coefficient: 16.9744
Is it a predictive feature?: Yes
Correlation: 0.3746637
Mean Square Error: 1109.453

Variable: tract_homevalue_2020
Intercept: 81.22493
Coefficient: -0.00003569445
Is it a predictive feature?: No
Correlation: 0.07536366
Mean Square Error: 1283.291

```

Variable: tract\_medage\_2020  
Intercept: -39.79579  
Coefficient: 2.790767  
Is it a predictive feature?: Yes  
Correlation: 0.4556463  
Mean Square Error: 1022.671

Variable: tract\_medincome\_2020  
Intercept: 41.09195  
Coefficient: 0.0003171854  
Is it a predictive feature?: Yes  
Correlation: 0.2672517  
Mean Square Error: 1198.441

Variable: tract\_medincome\_2010  
Intercept: 37.44741  
Coefficient: 0.0004233989  
Is it a predictive feature?: Yes  
Correlation: 0.3053277  
Mean Square Error: 1170.304

Variable: foreclosure\_pc\_2010  
Intercept: 20.68089  
Coefficient: 3.992378  
Is it a predictive feature?: Yes  
Correlation: 0.6491367  
Mean Square Error: 746.7816

Variable: pct\_built\_2020\_later  
Intercept: 68.74373  
Coefficient: 360.2854  
Is it a predictive feature?: No  
Correlation: 0.1201536  
Mean Square Error: 1271.989

Variable: pct\_built\_2010\_2019  
Intercept: 69.72607  
Coefficient: 9.976878  
Is it a predictive feature?: No  
Correlation: 0.02755463  
Mean Square Error: 1289.642

Variable: pct\_built\_2000\_2009

Intercept: 62.20988  
Coefficient: 83.76972  
Is it a predictive feature?: Yes  
Correlation: 0.2605284  
Mean Square Error: 1203.021

Variable: pct\_built\_1990\_1999  
Intercept: 58.84653  
Coefficient: 73.53769  
Is it a predictive feature?: Yes  
Correlation: 0.2630944  
Mean Square Error: 1201.287

Variable: pct\_built\_1980\_1989  
Intercept: 63.33133  
Coefficient: 48.99304  
Is it a predictive feature?: No  
Correlation: 0.146484  
Mean Square Error: 1262.928

Variable: pct\_built\_1970\_1979  
Intercept: 70.38708  
Coefficient: -0.9380611  
Is it a predictive feature?: No  
Correlation: 0.002405406  
Mean Square Error: 1290.614

Variable: pct\_built\_pre\_1960  
Intercept: 89.31805  
Coefficient: -46.31767  
Is it a predictive feature?: Yes  
Correlation: 0.326355  
Mean Square Error: 1153.161

Variable: pct\_0\_bd  
Intercept: 80.02659  
Coefficient: -482.1701  
Is it a predictive feature?: Yes  
Correlation: 0.3962199  
Mean Square Error: 1088.007

Variable: pct\_1\_bd  
Intercept: 85.44572

Coefficient: -133.0005  
Is it a predictive feature?: Yes  
Correlation: 0.4482354  
Mean Square Error: 1031.317

Variable: pct\_2\_bd  
Intercept: 87.1228  
Coefficient: -85.43262  
Is it a predictive feature?: Yes  
Correlation: 0.3832088  
Mean Square Error: 1101.095

Variable: pct\_3\_bd  
Intercept: 33.78992  
Coefficient: 116.2271  
Is it a predictive feature?: Yes  
Correlation: 0.422721  
Mean Square Error: 1059.997

Variable: pct\_4\_more\_bd  
Intercept: 53.82052  
Coefficient: 46.38325  
Is it a predictive feature?: Yes  
Correlation: 0.305918  
Mean Square Error: 1169.838

Variable: poverty\_2010  
Intercept: 87.40257  
Coefficient: -254.5372  
Is it a predictive feature?: Yes  
Correlation: 0.3931962  
Mean Square Error: 1091.087

Variable: poverty\_2020  
Intercept: 90.8764  
Coefficient: -265.7121  
Is it a predictive feature?: Yes  
Correlation: 0.3723224  
Mean Square Error: 1111.711

Variable: nhwhite\_2010  
Intercept: 78.41266  
Coefficient: -54.77476



Is it a predictive feature?: Yes  
Correlation: 0.2512609  
Mean Square Error: 1209.142

Variable: nhwhite\_2020  
Intercept: 77.66176  
Coefficient: -64.05339  
Is it a predictive feature?: Yes  
Correlation: 0.2492505  
Mean Square Error: 1210.441

Variable: mortgaged\_2010  
Intercept: 20.78945  
Coefficient: 93.89575  
Is it a predictive feature?: Yes  
Correlation: 0.607003  
Mean Square Error: 815.0888

Variable: mortgaged\_2015  
Intercept: 24.32348  
Coefficient: 94.48867  
Is it a predictive feature?: Yes  
Correlation: 0.5803045  
Mean Square Error: 856.0007

Variable: mortgaged\_2020  
Intercept: 24.26386  
Coefficient: 93.65878  
Is it a predictive feature?: Yes  
Correlation: 0.5779621  
Mean Square Error: 859.5023

Variable: ownoccupied\_2010  
Intercept: 22.55678  
Coefficient: 77.30333  
Is it a predictive feature?: Yes  
Correlation: 0.5639822  
Mean Square Error: 880.106

Variable: ownoccupied\_2015  
Intercept: 24.97609  
Coefficient: 77.30393  
Is it a predictive feature?: Yes

Correlation: 0.5406337  
Mean Square Error: 913.3927

Variable: ownoccupied\_2020  
Intercept: 24.14714  
Coefficient: 76.14728  
Is it a predictive feature?: Yes  
Correlation: 0.5469823  
Mean Square Error: 904.4811

Variable: mortgage\_change\_2010\_2015  
Intercept: 66.7023  
Coefficient: -87.32843  
Is it a predictive feature?: Yes  
Correlation: 0.1773345  
Mean Square Error: 1250.035

Variable: mortgage\_change\_2015\_2020  
Intercept: 70.2538  
Coefficient: 0.6904628  
Is it a predictive feature?: No  
Correlation: 0.001342858  
Mean Square Error: 1290.619

Variable: mortgage\_change\_2010\_2020  
Intercept: 68.15521  
Coefficient: 58.77525  
Is it a predictive feature?: No  
Correlation: 0.1449546  
Mean Square Error: 1263.503

Variable: ownoccupied\_change\_2010\_2015  
Intercept: 67.35991  
Coefficient: -92.56352  
Is it a predictive feature?: No  
Correlation: 0.175742  
Mean Square Error: 1250.761

Variable: ownoccupied\_change\_2015\_2020  
Intercept: 69.34088  
Coefficient: 46.3183  
Is it a predictive feature?: No  
Correlation: 0.08282612

Mean Square Error: 1281.768

Variable: ownoccupied\_change\_2010\_2020

Intercept: 69.90148

Coefficient: -30.88702

Is it a predictive feature?: No

Correlation: 0.0757052

Mean Square Error: 1283.225

Variable: poverty\_change\_2010\_2020

Intercept: 69.74779

Coefficient: 49.74486

Is it a predictive feature?: No

Correlation: 0.0652858

Mean Square Error: 1285.121

Variable: nhwhite\_change\_2010\_2020

Intercept: 72.14002

Coefficient: 56.55597

Is it a predictive feature?: No

Correlation: 0.1016417

Mean Square Error: 1277.288

Variable: medincome\_change\_2010\_2015

Intercept: 70.38322

Coefficient: -2.950833

Is it a predictive feature?: No

Correlation: 0.01251915

Mean Square Error: 1290.42

Variable: medincome\_change\_2015\_2020

Intercept: 74.16712

Coefficient: -24.11258

Is it a predictive feature?: No

Correlation: 0.1140783

Mean Square Error: 1273.826

Variable: medincome\_change\_2010\_2020

Intercept: 73.96636

Coefficient: -18.26635





Is it a predictive feature?: No

Correlation: 0.1064072

Mean Square Error: 1276.009

Variable: pop\_change\_pct  
Intercept: 71.01631  
Coefficient: -9.99906  
Is it a predictive feature?: No  
Correlation: 0.06121883  
Mean Square Error: 1285.785

We can also plot the relationships between our individual independent variables and our response variables using scatter plots. For example:



***Which are the most predictive features according to the training data?***

Evidently, the most predictive features are:

1. mortgaged\_2010/mortgaged\_2015/mortgaged\_2020. These are perhaps a little too obvious -- zip codes with more mortgaged homes are intuitively more likely to have a higher number of foreclosures per capita.

2. foreclosure\_pc\_2010. Once again, this is intuitive -- tracts that saw more foreclosures per capita during the height of the great recession likely have risk factors (included in our list of variables or not) that remained after 2010.

3. ownoccupied\_2010/ownoccupied\_2015/ownoccupied\_2020. The reasons for this relationship may be similar to the reasons for the relationship between the percentage of units with a mortgage and the per capita rate of foreclosure.

4. tract\_medage\_2020. According to the linear model, tracts with a higher median age are more likely to see higher rates of foreclosure per capita.

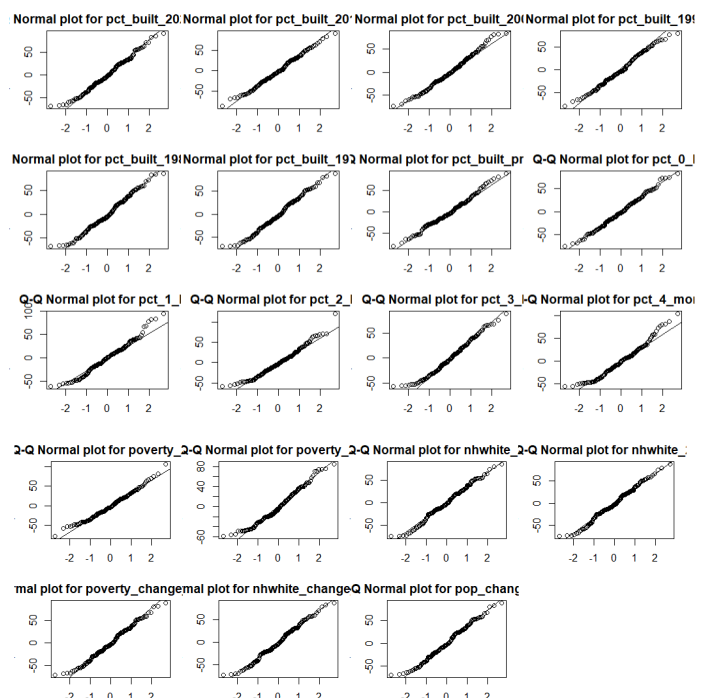
5. pct\_0\_bd/pct\_1\_bd/pct\_2\_bd/etc. The nature of the statistical relationship between the percentage of units in a tract with a given number of beds and the rate of foreclosures per capita is both difficult to explain and difficult to describe.

6. poverty\_2010. The tract-level poverty rate in 2020 is more predictive than the poverty rate in 2010 and the median income at the tract level in 2010 or 2020. This may be misleading -- a tract that saw a high rate of foreclosures between 2010-2020 may be poorer as a result of those foreclosures.

### ***What are the residuals? Is a linear regression applicable to your problem?***

Broadly speaking, the residuals for almost all variables follow a normal distribution, meaning a core assumption of linear regression (that residuals are normally distributed) are met.

That indicates that linear regression is generally applicable to our problem, though some variables – including pct\_1\_bd – may merit further investigation.



**Use the trained model to predict using your testing data. Show results together with confidence and prediction bands. Report prediction accuracy using (1) the correlation between the predicted and real values and (2) the mean square error between the two.**

Because of the number of independent variables within our dataset, we divided the variables of interest into subsets and wrote following code to loop through each variable to create a linear regression model, predict using the testing data, calculate the correlation and mean squared error, and plot the results:

```
for (variable in variables_of_interest_subset_1) {  
  # Perform Linear regression  
  model <- lm(foreclosure_pc_2020 ~ ., data = pg_foreclosures_train[,  
c(variable, "foreclosure_pc_2020")])  
  
  # Use the trained model to predict on testing data and obtain confidence  
  intervals  
  predictions <- predict(model, newdata = pg_foreclosures_test, interval =  
"confidence", level = 0.95)  
  
  # Calculate correlation between predicted and actual values  
  correlation_test <- cor(predictions[, "fit"],  
pg_foreclosures_test$foreclosure_pc_2020)  
  
  # Calculate mean square error  
  mse_test <- mean((predictions[, "fit"] -  
pg_foreclosures_test$foreclosure_pc_2020)^2)  
  
  # Plot predicted vs. actual values  
  plot(pg_foreclosures_test$foreclosure_pc_2020, predictions[, "fit"],  
    main = paste("Predicted vs. Actual for", variable),  
    xlab = "Actual Values", ylab = "Predicted Values", ylim =  
range(c(predictions[, "fit"], pg_foreclosures_test$foreclosure_pc_2020)))  
  abline(0, 1, col = "red") # Add a diagonal line for reference  
  
  # Print correlation and mean square error  
  cat("Variable:", variable, "\n")  
  cat("Correlation on Testing Data:", correlation_test, "\n")  
  cat("Mean Square Error on Testing Data:", mse_test, "\n")  
  cat("\n")  
}
```

We repeat the same code for each subset of independent variables and get the following results:

Variable: avg\_bed  
Correlation on Testing Data: 0.5117439  
Mean Square Error on Testing Data: 850.9713

Variable: tract\_homevalue\_2020  
Correlation on Testing Data: -0.3193279  
Mean Square Error on Testing Data: 1170.482

Variable: tract\_medage\_2020  
Correlation on Testing Data: 0.6158819  
Mean Square Error on Testing Data: 717.7623

Variable: tract\_medincome\_2020  
Correlation on Testing Data: 0.3378183  
Mean Square Error on Testing Data: 999.4049

Variable: tract\_medincome\_2010  
Correlation on Testing Data: 0.3635291  
Mean Square Error on Testing Data: 979.4372

Variable: medincome\_change\_2010\_2015  
Correlation on Testing Data: 0.08832201  
Mean Square Error on Testing Data: 1129.182

Variable: medincome\_change\_2015\_2020  
Correlation on Testing Data: 0.1959899  
Mean Square Error on Testing Data: 1091.314

Variable: medincome\_change\_2010\_2020  
Correlation on Testing Data: 0.1258555  
Mean Square Error on Testing Data: 1120.687

Variable: foreclosure\_pc\_2010  
Correlation on Testing Data: 0.3411653  
Mean Square Error on Testing Data: 1104.988

Variable: pct\_built\_2020\_later  
Correlation on Testing Data: 0.05592942  
Mean Square Error on Testing Data: 1128.079

Variable: pct\_built\_2010\_2019  
Correlation on Testing Data: -0.1423304  
Mean Square Error on Testing Data: 1191.664

Variable: pct\_built\_2000\_2009  
Correlation on Testing Data: 0.1736346  
Mean Square Error on Testing Data: 1140.387

Variable: pct\_built\_1990\_1999  
Correlation on Testing Data: 0.2315283  
Mean Square Error on Testing Data: 1073.282

Variable: pct\_built\_1980\_1989  
Correlation on Testing Data: 0.1130736  
Mean Square Error on Testing Data: 1117.865

Variable: pct\_built\_1970\_1979  
Correlation on Testing Data: 0.05390233  
Mean Square Error on Testing Data: 1132.112

Variable: pct\_built\_pre\_1960  
Correlation on Testing Data: 0.1382384  
Mean Square Error on Testing Data: 1211.898

Variable: pct\_0\_bd  
Correlation on Testing Data: 0.5573989  
Mean Square Error on Testing Data: 778.2076

Variable: pct\_1\_bd  
Correlation on Testing Data: 0.476484  
Mean Square Error on Testing Data: 876.044

Variable: pct\_2\_bd  
Correlation on Testing Data: 0.4531857  
Mean Square Error on Testing Data: 901.0094

Variable: pct\_3\_bd  
Correlation on Testing Data: 0.3998965  
Mean Square Error on Testing Data: 949.6106

Variable: pct\_4\_more\_bd  
Correlation on Testing Data: 0.4270028  
Mean Square Error on Testing Data: 930.5358

Variable: poverty\_2010  
Correlation on Testing Data: 0.5019256  
Mean Square Error on Testing Data: 843.8109

Variable: poverty\_2020  
Correlation on Testing Data: 0.05570013  
Mean Square Error on Testing Data: 1214

Variable: nhwhite\_2010  
Correlation on Testing Data: 0.3283699



Mean Square Error on Testing Data: 1059.075

Variable: nhwhite\_2020

Correlation on Testing Data: 0.4733185

Mean Square Error on Testing Data: 1049.925

Variable: poverty\_change\_2010\_2020

Correlation on Testing Data: 0.5307815

Mean Square Error on Testing Data: 1096.449

Variable: nhwhite\_change\_2010\_2020

Correlation on Testing Data: 0.05467011

Mean Square Error on Testing Data: 1128.522

Variable: pop\_change\_pct

Correlation on Testing Data: 0.1093619

Mean Square Error on Testing Data: 1120.843

Variable: mortgaged\_2010

Correlation on Testing Data: 0.6191706

Mean Square Error on Testing Data: 699.3978

Variable: mortgaged\_2015

Correlation on Testing Data: 0.5824819

Mean Square Error on Testing Data: 768.5658

Variable: mortgaged\_2020

Correlation on Testing Data: 0.5966077

Mean Square Error on Testing Data: 735.8966

Variable: ownoccupied\_2010

Correlation on Testing Data: 0.6195836

Mean Square Error on Testing Data: 695.9821

Variable: ownoccupied\_2015

Correlation on Testing Data: 0.5838292

Mean Square Error on Testing Data: 744.8564

Variable: ownoccupied\_2020

Correlation on Testing Data: 0.5924652

Mean Square Error on Testing Data: 735.3165

Variable: mortgage\_change\_2010\_2015

Correlation on Testing Data: 0.06476888

Mean Square Error on Testing Data: 1146.255

Variable: mortgage\_change\_2015\_2020

Correlation on Testing Data:  $-0.04575922$   
Mean Square Error on Testing Data:  $1142.707$

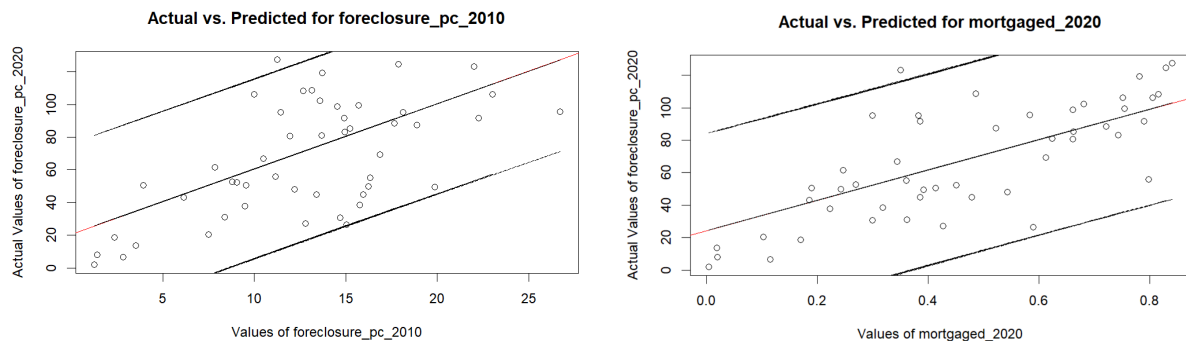
Variable: mortgage\_change\_2010\_2020  
Correlation on Testing Data:  $0.09506252$   
Mean Square Error on Testing Data:  $1123.147$

Variable: ownoccupied\_change\_2010\_2015  
Correlation on Testing Data:  $0.1337618$   
Mean Square Error on Testing Data:  $1110.936$

Variable: ownoccupied\_change\_2015\_2020  
Correlation on Testing Data:  $0.1337533$   
Mean Square Error on Testing Data:  $1110.026$

Variable: ownoccupied\_change\_2010\_2020  
Correlation on Testing Data:  $0.005473917$   
Mean Square Error on Testing Data:  $1134.188$

It would be too space-intensive to include plots for each of the 40 variables, but here is sample:



We did not successfully plot both confidence and prediction bands; we struggled to apply the `matlines()` function to the loop we used to train, test and plot a linear regression model for each independent variable.

## B. Multivariate regressions

**Show whether considering combinations of independent features improves the prediction results. Evaluate different combinations of features as applicable and report those that improve the results shown in question (a).**

As a starting point, we use the following code to identify collinear groups of variables using a threshold correlation coefficient of 0.8:

```

# Calculating the correlation matrix
correlation_matrix <- cor(pg_foreclosures_per_tract[variables_of_interest])

# Set a correlation threshold
threshold <- 0.85

# Find highly correlated variable pairs
highly_correlated <- which(correlation_matrix > threshold &
correlation_matrix < 1, arr.ind = TRUE)

# Print highly correlated variable pairs
for (i in 1:nrow(highly_correlated)) {
  var1 <- rownames(correlation_matrix)[highly_correlated[i, 1]]
  var2 <- colnames(correlation_matrix)[highly_correlated[i, 2]]
  corr <- correlation_matrix[highly_correlated[i, 1], highly_correlated[i,
2]]
  cat("Variables", var1, "and", var2, "are highly correlated (correlation
=", corr, ")\n")
}

```

For our first experimental set of independent variables, we select the most predictive independent variable from each group of collinear variables, add any variables that were not collinear with other variables, and create the following subset:

```

test_variables_1 <- c("avg_bed", "tract_homevalue_2020",
"tract_medage_2020", "tract_medincome_2020", "foreclosure_pc_2010",
"pct_built_2020_later", "pct_built_2010_2019", "pct_built_2000_2009",
"pct_built_1990_1999", "pct_built_1980_1989", "pct_built_1970_1979",
"pct_built_pre_1960", "poverty_2020", "nhwhite_2020", "mortgaged_2010",
"mortgage_change_2010_2015",
"poverty_change_2010_2020", "nhwhite_change_2010_2020", "medincome_change_201
0_2015", "medincome_change_2015_2020", "medincome_change_2010_2020",
"pop_change_pct")

```

We then run the following code to divide our dataset into randomized training and test sets, train the linear regression model using the selected variables, use the trained model to predict the response variable, and calculate the **correlation between real and predicted values, the mean squared error between the real and predicted values, the coefficients for each feature, and the most predictive features of the model**:

```

# Convert tract_number to character before splitting the data

```

```

pg_foreclosures_per_tract$tract_number <-
as.character(pg_foreclosures_per_tract$tract_number)

# Split the data into training and testing sets
set.seed(123) # for reproducibility
train_indices <-
createDataPartition(pg_foreclosures_per_tract$foreclosure_pc_2020, p = 0.8,
list = FALSE)
train_data_multivariate_1 <- pg_foreclosures_per_tract[train_indices, ]
test_data_multivariate_1 <- pg_foreclosures_per_tract[-train_indices, ]

# Train the linear regression model on the training data
lm_model <- lm(foreclosure_pc_2020 ~ ., data = train_data_multivariate_1[,
c("foreclosure_pc_2020", test_variables_1)])

# Use the trained model to predict on the testing data
predicted_values <- predict(lm_model, newdata = test_data_multivariate_1[,
test_variables_1])

# Calculate the correlation between the predicted and real values
correlation <- cor(predicted_values,
test_data_multivariate_1$foreclosure_pc_2020)

# Calculate the mean squared error between the predicted and real values
mse <- mean((predicted_values -
test_data_multivariate_1$foreclosure_pc_2020)^2)

# Obtain coefficients for each feature
coefficients <- coef(lm_model)
print(coefficients)

# Identify the most predictive features
# Absolute values of coefficients can be considered for importance
# Higher absolute values indicate more influence on the prediction
absolute_coefficients <- abs(coefficients[-1]) # Exclude intercept
top_predictive_features <-
names(absolute_coefficients)[order(-absolute_coefficients)][1:5] # Select
top 5 features

```

That initial model produces the following output:

Coefficients:

```
avg_bed: -3.270173433910
tract_homevalue_2020: -0.000071088618
tract_medage_2020: 0.943090883386
tract_medincome_2020: -0.000007804235
foreclosure_pc_2010: 2.600806436706
pct_built_2020_later: 209.729556742509
pct_built_2010_2019: 78.823749017584
pct_built_2000_2009: 19.867570073357
pct_built_1990_1999: 41.230919498650
pct_built_1980_1989: -35.810641773229
pct_built_1970_1979: 31.290999117118
pct_built_pre_1960: NA
poverty_2020: -107.531979132470
nhwhite_2020: -62.911825462744
mortgaged_2010: 62.545986109374
mortgage_change_2010_2015: 40.321693073811
poverty_change_2010_2020: 58.644184450053
nhwhite_change_2010_2020: 40.137487023411
medincome_change_2010_2015: -108.743283778538
medincome_change_2015_2020: -112.855022494620
medincome_change_2010_2020: 102.354402037439
pop_change_pct: -42.460893747868
```

Most predictive features:

```
pct_built_2020_later
medincome_change_2015_2020
medincome_change_2010_2015
poverty_2020
medincome_change_2010_2020
```

Correlation between predicted and real values: 0.7878868

Mean Squared Error: 434.8585

That model did, however, also produce this warning, which indicated some degree of collinearity:

```
Warning: prediction from a rank-deficient fit may be misleading
```

After some experimentation, we reached this set of independent variables that both improved the predictiveness of our model and did not produce a warning about rank-deficient fit:

```
test_variables_3 <- c("mortgaged_2010", "tract_homevalue_2020",  
"tract_medage_2020", "tract_medincome_2020", "foreclosure_pc_2010",  
"pct_built_2000_2009", "poverty_2010", "nhwhite_2020",  
"mortgage_change_2010_2020",  
"poverty_change_2010_2020", "nhwhite_change_2010_2020", "medincome_change_2010_2015",  
"medincome_change_2015_2020", "medincome_change_2010_2020",  
"pop_change_pct")
```

That combination of variables produced this output:

Coefficients:

```
mortgaged_2010: 22.797472442326  
tract_homevalue_2020: 62.996602305439  
tract_medage_2020: -0.000074835506  
tract_medincome_2020: -0.000004506972  
foreclosure_pc_2010: 2.235976962519  
pct_built_2000_2009: 47.156334551554  
poverty_2010: -126.862854350235  
nhwhite_2020: -71.672688012288  
mortgage_change_2010_2020: -64.941360273727  
poverty_change_2010_2020: -70.380936131425  
nhwhite_change_2010_2020: 60.125954488660  
medincome_change_2010_2015: -143.305801287334  
medincome_change_2015_2020: -140.909399442653  
medincome_change_2010_2020: 128.456180060477  
pop_change_pct: -27.094446208530
```

Most predictive features:

```
medincome_change_2010_2015  
medincome_change_2015_2020  
medincome_change_2010_2020  
poverty_2010  
nhwhite_2020
```

Correlation between predicted and real values: 0.7934256

Mean Squared Error: 437.2731

Of the models we tested, that model was the most successful.

### C. Regularization

**Repeat experiments in (a) and (b) adding regularization. Do you observe any improvements in the prediction results?**

We used the following code to add regularization to our calculation of univariate logistic regressions for each variable:

```
# Set the regularization parameter
lambda <- 0.2

for (variable in variables_of_interest) {
  # Prepare the data
  X <- model.matrix(foreclosure_pc_2020 ~ ., data =
pg_foreclosures_train[, c(variables_of_interest, "foreclosure_pc_2020")][,
-1])
  y <- pg_foreclosures_train$foreclosure_pc_2020

  # Fit the Lasso regression model
  lasso_model <- glmnet(X, y, alpha = 1, lambda = lambda)

  # Extract coefficients
  coef_idx <- which(colnames(X) == variable)
  coefficient <- coef(lasso_model)[coef_idx]

  # Check if the coefficient is non-zero
  is_predictive <- ifelse(abs(coefficient) > 0, "Yes", "No")

  # Compute residuals
  residuals <- y - predict(lasso_model, newx = X)

  # Compute correlation between predicted and real values
  correlation <- cor(predict(lasso_model, newx = X), y)

  # Compute mean square error
  mse <- mean(residuals^2)

  # Print results
  cat("Variable:", variable, "\n")
  cat("Coefficient:", coefficient, "\n")
  cat("Is it a predictive feature?:", is_predictive, "\n")
  cat("Correlation:", correlation, "\n")
  cat("Mean Square Error:", mse, "\n")
  cat("\n")
}
```

Regularization produces consistent correlation coefficients of roughly 0.87 for each variable – an improvement – though the coefficients for each variable remain distinct.

After regularization, the residuals remain normally distributed.

We used the following code to use the regularized model on our test data:

```
# Convert tract_number to character before splitting the data
pg_foreclosures_per_tract$tract_number <-
as.character(pg_foreclosures_per_tract$tract_number)

# Split the data into training and testing sets
set.seed(123) # for reproducibility
train_indices <-
createDataPartition(pg_foreclosures_per_tract$foreclosure_pc_2020, p = 0.8,
list = FALSE)
train_data_multivariate_1 <- pg_foreclosures_per_tract[train_indices, ]
test_data_multivariate_1 <- pg_foreclosures_per_tract[-train_indices, ]

# Define predictor variables
test_variables_3 <- c("ownoccupied_2010", "tract_homevalue_2020",
"tract_medage_2020", "tract_medincome_2020", "foreclosure_pc_2010",
"pct_built_2000_2009", "poverty_2010", "nhwhite_2020",
"mortgage_change_2010_2020",
"poverty_change_2010_2020", "nhwhite_change_2010_2020", "medincome_change_2010_2015", "medincome_change_2015_2020", "medincome_change_2010_2020",
"pop_change_pct")

# Train the Lasso regression model on the training data
lasso_model <- cv.glmnet(as.matrix(train_data_multivariate_1[,
test_variables_3]),
train_data_multivariate_1$foreclosure_pc_2020,
alpha = 1) # Use alpha = 1 for Lasso regression

# Use the trained model to predict on the testing data
predicted_values <- predict(lasso_model, newx =
as.matrix(test_data_multivariate_1[, test_variables_3]),
s = "lambda.min")
```



```

# Calculate the correlation between the predicted and real values
correlation <- cor(predicted_values,
test_data_multivariate_1$foreclosure_pc_2020)

# Calculate the mean squared error between the predicted and real values
mse <- mean((predicted_values -
test_data_multivariate_1$foreclosure_pc_2020)^2)

# Obtain coefficients for each feature
coefficients <- coef(lasso_model, s = "lambda.min")
print(coefficients)

# Identify the most predictive features
# Non-zero coefficients indicate predictive features in Lasso regression
non_zero_coefficients <- coefficients[-1] # Exclude intercept
top_predictive_features <-
names(non_zero_coefficients[non_zero_coefficients != 0]) # Select features
with non-zero coefficients

# Print the most predictive features, correlation, and mean squared error
cat("Most predictive features:", top_predictive_features, "\n")
cat("Correlation between predicted and real values:", correlation, "\n")
cat("Mean Squared Error:", mse, "\n")

```

That code produces this output:

```

(Intercept)          5.86204400487
ownoccupied_2010      53.36979029348
tract_homevalue_2020  -0.00005998712
tract_medage_2020     0.92702987857
tract_medincome_2020  .
foreclosure_pc_2010   2.22242203802
pct_built_2000_2009   48.93320838235
poverty_2010          -84.76763233288
nhwhite_2020          -74.74020982294
mortgage_change_2010_2020 -45.66978992889
poverty_change_2010_2020 -36.48885980937
nhwhite_change_2010_2020  54.28663831913
medincome_change_2010_2015  0.79485266353
medincome_change_2015_2020 -0.99014238712
medincome_change_2010_2020 .
pop_change_pct        -22.56302235567
Most predictive features:
Correlation between predicted and real values: 0.7758073

```

Mean Squared Error: 456.9231

After using the regularized trained model on our test data, we found no "most predictive" features. In practical terms, we can interpret this result as the Lasso algorithm finding that no single feature is dominant in predicting the foreclosure percentages, and the model relies on a combination of features with small contributions from each. This is consistent with the identical correlation coefficients calculated for each of our independent variables using the regularized univariate regression model.

Unfortunately, the regularized model is slightly less accurate than the initial multivariate regression model.

***D. Repeat a-c multiple times with different randomly selected training and testing sets and report differences or similarities across runs.***

We used the following code to repeat the initial univariate regression model process with multiple unique training and test sets:

```
# Set seed for reproducibility
set.seed(123)

# Define the number of repetitions
num_repetitions <- 5

# Loop for repetitions
for (i in 1:num_repetitions) {
  # Split data into training and testing datasets (e.g., 80% training, 20%
  # testing)
  train_indices <-
  createDataPartition(pg_foreclosures_per_tract$foreclosure_pc_2020, p = 0.7,
  list = FALSE)
  pg_foreclosures_train <- pg_foreclosures_per_tract[train_indices, ]
  pg_foreclosures_test <- pg_foreclosures_per_tract[-train_indices, ]

  for (variable in variables_of_interest) {
    # Perform linear regression
    model <- lm(foreclosure_pc_2020 ~ ., data = pg_foreclosures_train[,
    c(variable, "foreclosure_pc_2020")])

    # Extract required information
    intercept <- coef(model)[1]
    coefficient <- coef(model)[2]
```

```

    # Check if the coefficient is statistically significant
    p_value <-
summary(model)$coefficients[which(rownames(summary(model)$coefficients) ==
variable), "Pr(>|t|)"]
    is_predictive <- ifelse(p_value < 0.05, "Yes", "No")

    # Compute residuals
    residuals <- resid(model)

    # Compute correlation between predicted and real values
    correlation <- cor(predict(model),
pg_foreclosures_train$foreclosure_pc_2020 )

    # Compute mean square error
    mse <- mean((predict(model) - pg_foreclosures_train$foreclosure_pc_2020
)^2)

    # Print results
    cat("Repetition:", i, "\n")
    cat("Variable:", variable, "\n")
    cat("Intercept:", intercept, "\n")
    cat("Coefficient:", coefficient, "\n")
    cat("Is it a predictive feature?:", is_predictive, "\n")
    cat("Correlation:", correlation, "\n")
    cat("Mean Square Error:", mse, "\n")
    cat("\n")
  }
}

```

And the following code to test our most successful multivariate regression model with multiple unique training and test sets:

```

# Set the number of iterations
num_iterations <- 5

# Initialize empty vectors to store results
correlations <- numeric(num_iterations)
mses <- numeric(num_iterations)

for (i in 1:num_iterations) {
  # Convert tract_number to character before splitting the data

```

```

pg_foreclosures_per_tract$tract_number <-
as.character(pg_foreclosures_per_tract$tract_number)

# Split the data into training and testing sets
set.seed(i) # Use different seed for each iteration
train_indices <-
createDataPartition(pg_foreclosures_per_tract$foreclosure_pc_2020, p = 0.8,
list = FALSE)
train_data <- pg_foreclosures_per_tract[train_indices, ]
test_data <- pg_foreclosures_per_tract[-train_indices, ]

# Train the linear regression model on the training data
lm_model <- lm(foreclosure_pc_2020 ~ ., data = train_data[,
c("foreclosure_pc_2020", test_variables_3)])

# Use the trained model to predict on the testing data
predicted_values <- predict(lm_model, newdata = test_data[,
test_variables_3])

# Calculate the correlation between the predicted and real values
correlation <- cor(predicted_values, test_data$foreclosure_pc_2020)

# Calculate the mean squared error between the predicted and real values
mse <- mean((predicted_values - test_data$foreclosure_pc_2020)^2)

# Store correlation and mse
correlations[i] <- correlation
mses[i] <- mse
}

# Print the results of each iteration
for (i in 1:num_iterations) {
  cat("Iteration", i, "\n")
  cat("Correlation between predicted and real values:", correlations[i],
"\n")
  cat("Mean Squared Error:", mses[i], "\n\n")
}

```

And the following code to repeat the regularized univariate regression model process for each of our independent variables:

```

# Set the regularization parameter
lambda <- 0.2

# Number of iterations
num_iterations <- 5

for (iteration in 1:num_iterations) {
  # Split data into training and test sets
  set.seed(iteration) # Set seed for reproducibility
  sample_indices <- sample(1:nrow(pg_foreclosures_per_tract), size = 0.7
* nrow(pg_foreclosures_per_tract), replace = FALSE)
  pg_foreclosures_train <- pg_foreclosures_per_tract[sample_indices, ]
  pg_foreclosures_test <- pg_foreclosures_per_tract[-sample_indices, ]

  for (variable in variables_of_interest) {
    # Prepare the data
    X_train <- model.matrix(foreclosure_pc_2020 ~ ., data =
pg_foreclosures_train[, c(variables_of_interest, "foreclosure_pc_2020")][,
-1])
    y_train <- pg_foreclosures_train$foreclosure_pc_2020
    X_test <- model.matrix(foreclosure_pc_2020 ~ ., data =
pg_foreclosures_test[, c(variables_of_interest, "foreclosure_pc_2020")][,
-1])
    y_test <- pg_foreclosures_test$foreclosure_pc_2020

    # Fit the Lasso regression model
    lasso_model <- glmnet(X_train, y_train, alpha = 1, lambda = lambda)

    # Extract coefficients
    coef_idx <- which(colnames(X_train) == variable)
    coefficient <- coef(lasso_model)[coef_idx]

    # Check if the coefficient is non-zero
    is_predictive <- ifelse(abs(coefficient) > 0, "Yes", "No")

    # Compute residuals
    residuals <- y_test - predict(lasso_model, newx = X_test)

    # Compute correlation between predicted and real values
    correlation <- cor(predict(lasso_model, newx = X_test), y_test)

    # Compute mean square error
    mse <- mean(residuals^2)
  }
}

```

```

    # Print results
    cat("Iteration:", iteration, "\n")
    cat("Variable:", variable, "\n")
    cat("Coefficient:", coefficient, "\n")
    cat("Is it a predictive feature?:", is_predictive, "\n")
    cat("Correlation:", correlation, "\n")
    cat("Mean Square Error:", mse, "\n")
    cat("\n")
  }
}

```

And the following code to test the most regularized multivariate regression model using multiple unique training and test sets:

```

# Set the number of repetitions
num_repetitions <- 5

# Loop through repetitions
for (i in 1:num_repetitions) {
  # Split the data into training and testing sets
  set.seed(i) # Use different seed for each iteration for
  reproducibility
  train_indices <-
  createDataPartition(pg_foreclosures_per_tract$foreclosure_pc_2020, p = 0.8,
  list = FALSE)
  train_data <- pg_foreclosures_per_tract[train_indices, ]
  test_data <- pg_foreclosures_per_tract[-train_indices, ]

  # Train the Lasso regression model on the training data
  lasso_model <- cv.glmnet(as.matrix(train_data[, test_variables_3]),
                           train_data$foreclosure_pc_2020,
                           alpha = 1)

  # Use the trained model to predict on the testing data
  predicted_values <- predict(lasso_model, newx = as.matrix(test_data[,
  test_variables_3]),
                             s = "lambda.min")

  # Calculate the correlation between the predicted and real values
  correlation <- cor(predicted_values, test_data$foreclosure_pc_2020)
}

```

```

# Calculate the mean squared error between the predicted and real
values
mse <- mean((predicted_values - test_data$foreclosure_pc_2020)^2)

# Obtain coefficients for each feature
coefficients <- coef(lasso_model, s = "lambda.min")

# Identify the most predictive features
non_zero_coefficients <- coefficients[-1] # Exclude intercept
top_predictive_features <-
names(non_zero_coefficients[non_zero_coefficients != 0])

# Print the results for this run
cat("Run:", i, "\n")
cat("Correlation between predicted and real values:", correlation,
"\n")
cat("Mean Squared Error:", mse, "\n")
cat("Top predictive features:", top_predictive_features, "\n")
cat("\n")
}

```

For the sake of space, we will only show the output from the final test, which was fairly representative of the other tests.

That code produces this output:

```

Run: 1
Correlation between predicted and real values: 0.8422236
Mean Squared Error: 385.4891
Top predictive features:

Run: 2
Correlation between predicted and real values: 0.8523196
Mean Squared Error: 384.6485
Top predictive features:

Run: 3
Correlation between predicted and real values: 0.8180499
Mean Squared Error: 384.1731
Top predictive features:

Run: 4

```

```
Correlation between predicted and real values: 0.8194143
Mean Squared Error: 340.0986
Top predictive features:
```

```
Run: 5
Correlation between predicted and real values: 0.8901317
Mean Squared Error: 293.5396
Top predictive features:
```

In other words, the predictiveness of our regularized multivariate regression model can vary substantially depending on the training and test sets used to build and test it.

The same is true for our unregularized multivariate regression model and both the unregularized and regularized univariate linear regression models, which could indicate that our relatively small sample size makes it difficult to create a model that does not vary significantly in its predictiveness depending on the sample used to train it.

## **Question 2. Logistic Regression and Naive Bayes:**

**Preparation:** In order to create a logistic regression model and Naive Bayes model for our dataset, we first needed to turn our numeric response variable, `foreclosure_pc_2020` – which represents the number of foreclosures in a tract between 2011-2023 divided by the most recent ACS population estimate for said tract – into an ordinal variable. We do this by dividing the range of `foreclosure_pc_2020` values into quantiles, each of which becomes a class within the ordinal variable `foreclosure_quantile`.

### **A. Logistic Regression:**

After that, we built a proportional odds logistic regression formula – which is more appropriate for predicting ordinal variables – to build our model.

For some reason, the exact combination of variables from our most predictive linear regression model did not run successfully when input into our logistic regression model.

For reference, here is that combination of variables:

```
test_variables_3 <- c("mortgaged_2010", "tract_homevalue_2020",
"tract_medage_2020", "tract_medincome_2020", "foreclosure_pc_2010",
"pct_built_2000_2009", "poverty_2010", "nhwhite_2020",
"mortgage_change_2010_2020",
"poverty_change_2010_2020", "nhwhite_change_2010_2020", "medincome_change_2010_2015", "medincome_change_2015_2020", "medincome_change_2010_2020",
"pop_change_pct")
```



Running that combination of variables through the logistic regression model looks like this:

```
train_data$foreclosure_quantile <- factor(train_data$foreclosure_quantile)

# Use the training dataset for model fitting
model_1 <- polr(foreclosure_quantile ~ mortgaged_2010 +
  tract_homevalue_2020 + tract_medage_2020 + tract_medincome_2020 +
  foreclosure_pc_2010 + pct_built_2000_2009 + poverty_2010 + nhwhite_2020 +
  mortgage_change_2010_2020 + poverty_change_2010_2020 +
  nhwhite_change_2010_2020 + medincome_change_2010_2015 +
  medincome_change_2015_2020 + medincome_change_2010_2020 + pop_change_pct,
  data = train_data, Hess = TRUE)

# Summarize the model
summary(model_1)
```

We get this error:

```
Error in svd(X) : infinite or missing values in 'x'
```

That error generally indicates missing or infinite values, but because there are neither infinite nor missing values in our dataset, it more likely means that there are collinear variables in our model.

In our linear regression section, we set the bar for collinear variables at a correlation coefficient (between independent variables) above 0.8.

It appears that our logistic regression model is more sensitive, so we ran a new test to correlation coefficients above 0.7 using this code:

```
# Select the variables from the regression formula
selected_vars <- c("mortgaged_2010", "tract_homevalue_2020",
  "tract_medage_2020",
  "tract_medincome_2020", "foreclosure_pc_2010",
  "pct_built_2000_2009",
  "poverty_2010", "nhwhite_2020",
  "mortgage_change_2010_2020",
  "poverty_change_2010_2020", "nhwhite_change_2010_2020",
  "medincome_change_2010_2015",
  "medincome_change_2015_2020",
```

```

    "medincome_change_2010_2020", "pop_change_pct")

# Subset the data with selected variables
selected_data <- train_data[selected_vars]

# Calculate the correlation matrix
correlation_matrix <- cor(selected_data)

# Find pairs of variables with correlation coefficient > 0.7
high_correlation <- which(upper.tri(correlation_matrix, diag = TRUE) &
correlation_matrix > 0.7, arr.ind = TRUE)

# Print the collinear variable pairs
collinear_pairs <- data.frame(variable1 =
rownames(correlation_matrix)[high_correlation[, 1]],
                             variable2 =
colnames(correlation_matrix)[high_correlation[, 2]],
                             correlation =
correlation_matrix[high_correlation])

```

The resulting correlation matrix pointed us to two potential pairs of problem variables:

mortgaged_2010	tract_medincome_2020	0.7636285
tract_homevalue_2020	tract_medincome_2020	0.7285017

Of those, mortgaged\_2010 was most statistically significant (as calculated during the linear regression section), so we jettisoned the substantially less-predictive variables tract\_medincome\_2020 and tract\_homevalue\_2020 and keep only mortgaged\_2010.

**Model 1:** That gave us the following model as a starting point:

```

# Use the training dataset for model fitting
model_1 <- polr(foreclosure_quantile ~ mortgaged_2010 + tract_medage_2020 +
foreclosure_pc_2010 + pct_built_2000_2009 + poverty_2010 + nhwhite_2020 +
mortgage_change_2010_2020 + poverty_change_2010_2020 +
nhwhite_change_2010_2020 + medincome_change_2010_2015 +
medincome_change_2015_2020 + medincome_change_2010_2020 + pop_change_pct,
data = train_data, Hess = TRUE)

```

The following are the results of the initial model, which provides both the **coefficients** for each variable and the **intercepts** for the transitions between each quantile:

Coefficients:

	Value	Std. Error	t value
mortgaged_2010	5.1645	1.40460	3.677
tract_medage_2020	0.1112	0.04456	2.495
foreclosure_pc_2010	0.1988	0.04297	4.627
pct_built_2000_2009	5.3423	2.34005	2.283
poverty_2010	-9.0151	5.84260	-1.543
nhwhite_2020	-8.4189	2.02452	-4.158
mortgage_change_2010_2020	-9.2956	2.78059	-3.343
poverty_change_2010_2020	-5.1738	4.77354	-1.084
nhwhite_change_2010_2020	5.3517	3.08483	1.735
medincome_change_2010_2015	-11.8941	6.80579	-1.748
medincome_change_2015_2020	-10.3685	6.24906	-1.659
medincome_change_2010_2020	9.1935	5.88229	1.563
pop_change_pct	-3.0799	0.94985	-3.243

Intercepts:

	Value	Std. Error	t value
1 2	4.3986	1.8510	2.3763
2 3	6.6132	1.9139	3.4554
3 4	8.6115	1.9879	4.3320
4 5	10.4258	2.0457	5.0963

Residual Deviance: 258.4844

AIC: 292.4844

### ***Are the coefficients statistically significant?***

We used the t-values to determine whether the coefficients were statistically significant. Based on the calculated t-values, the coefficients for foreclosure\_pc\_2010 (t-value = 4.627), mortgaged\_2010 (t-value = 3.677), and nhwhite\_2020 (t-value = -4.158) are the most statistically significant.

### ***What are the log-odds and odd ratios of the outcome for a unit increase in each independent variable?***

As far as we understand, the [coefficients in logistic regression represent the change in the log-odds of the outcome variable for a one-unit change in the predictor variable](#). Under that interpretation, the log-odds should be equivalent to the coefficients for each variable.

Under that interpretation, the following are the **odds ratios** for a unit increase in each independent variable:

1. mortgaged\_2010: 174.947
2. tract\_medage\_2020: 1.118
3. foreclosure\_pc\_2010: 1.220
4. pct\_built\_2000\_2009: 209.996
5. poverty\_2010: 0.0001216
6. nhwhite\_2020: 0.0002207
7. mortgage\_change\_2010\_2020: 0.00009183
8. poverty\_change\_2010\_2020: 0.005663
9. nhwhite\_change\_2010\_2020: 210.964
10. medincome\_change\_2010\_2015: 0.000006830
11. medincome\_change\_2015\_2020: 0.00003141
12. medincome\_change\_2010\_2020: 9832.97
13. pop\_change\_pct: 0.04596

***Which are the most predictive features according to the training data?***

We rank the predictiveness of each variable based on both the significance of the variable (indicated by the t-value) and the effect size (indicated by the log-odds). Variables with larger absolute log-odds and t-values are considered **most predictive** in our ranking.

1. foreclosure\_pc\_2010 (t-value: 4.627, log-odds: 0.1988)
2. nhwhite\_2020 (t-value: -4.158, log-odds: -8.4189)
3. mortgaged\_2010 (t-value: 3.677, log-odds: 5.1645)
4. mortgage\_change\_2010\_2020 (t-value: -3.343, log-odds: -9.2956)
5. pop\_change\_pct (t-value: -3.243, log-odds: -3.0799)
6. tract\_medage\_2020 (t-value: 2.495, log-odds: 0.1112)
7. pct\_built\_2000\_2009 (t-value: 2.283, log-odds: 5.3423)
8. nhwhite\_change\_2010\_2020 (t-value: 1.735, log-odds: 5.3517)
9. medincome\_change\_2010\_2020 (t-value: 1.563, log-odds: 9.1935)
10. medincome\_change\_2015\_2020 (t-value: -1.659, log-odds: -10.3685)
11. medincome\_change\_2010\_2015 (t-value: -1.748, log-odds: -11.8941)
12. poverty\_change\_2010\_2020 (t-value: -1.084, log-odds: -5.1738)
13. poverty\_2010 (t-value: -1.543, log-odds: -9.0151)

As expected, foreclosure\_pc\_2010 is still the most predictive variable; unexpectedly, nhwhite\_2020 is the second-most predictive.

***Use the trained model to predict on your testing dataset.***

We used the following code to predict foreclosure\_quantile values based on the values in a test set:

```
# Predict values using the model and test data
predicted_values <- predict(model_1, newdata = test_data, type = "class")


# Calculate the correlation between predicted and actual values
correlation <- cor(as.numeric(predicted_values),
test_data$foreclosure_quantile)
```

The resulting correlation coefficient between our predicted results and actual values is **0.82**.

For the sake of thoroughness, we also conducted a one-to-one comparison of the predicted quantiles and actual quantiles to calculate an accuracy rate.

This model accurately predicted quantiles roughly 54% of the time.

That outcome was far from ideal, so **we tested other combinations of variables in search of a more predictive model.**



**Model 2:** In the training phase, the following model was the second-most predictive:

```
# Use the training dataset for model fitting
model_6 <- polr(foreclosure_quantile ~ foreclosure_pc_2010 +
tract_medage_2020 + poverty_2010 + nhwhite_2020 +
medincome_change_2010_2020 + avg_bed + mortgaged_2010 + pct_1_bd +
pct_built_pre_1960 + pct_built_2000_2009 + mortgage_change_2010_2015 +
pop_change_pct, data = train_data, Hess = TRUE)

# Summarize the model
summary(model_6)
```

That model produced the following output, at least when considering the residual deviance and AIC:

```
Coefficients:
Value Std. Error t value
foreclosure_pc_2010 0.1455 0.04298 3.3848
```

tract_medage_2020	0.1587	0.04780	3.3207
poverty_2010	-2.1505	4.47238	-0.4808
nhwhite_2020	-8.6907	2.00621	-4.3319
medincome_change_2010_2020	1.1254	0.94781	1.1874
avg_bed	-1.3916	0.57201	-2.4329
mortgaged_2010	5.7754	2.57321	2.2444
pct_1_bd	-7.0085	3.31349	-2.1151
pct_built_pre_1960	-0.8271	0.93078	-0.8886
pct_built_2000_2009	3.4087	2.50129	1.3628
mortgage_change_2010_2015	2.7022	3.11060	0.8687
pop_change_pct	-2.8016	0.95431	-2.9357

Intercepts:

	Value	Std. Error	t value
1 2	2.1020	2.0460	1.0274
2 3	4.2130	2.0839	2.0217
3 4	5.9874	2.1305	2.8103
4 5	7.6469	2.1608	3.5390

Residual Deviance: 269.425

AIC: 301.425

### ***Are the coefficients statistically significant?***

Based on the calculated t-values, the coefficients for foreclosure\_pc\_2010 (t-value = 5.573), tract\_medage\_2020 (t-value = 3.894), nhwhite\_2020 (t-value = -4.3319), and mortgaged\_2010 (t-value = 3.677) are the are statistically significant.

### ***What are the odd ratios of the outcome for a unit increase in each independent variable?***

The **odds ratios** for a unit increase in each independent variable with this model are:

1. foreclosure\_pc\_2010: 1.1565860422
2. tract\_medage\_2020: 1.1720011361
3. poverty\_2010: 0.1164288069
4. nhwhite\_2020: 0.0001681451
5. medincome\_change\_2010\_2020: 3.0815814210
6. avg\_bed: 0.2486755107
7. mortgaged\_2010: 322.2597426584
8. pct\_1\_bd: 0.0009041991
9. pct\_built\_pre\_1960: 0.4373152291
10. pct\_built\_2000\_2009: 30.2260067268

### ***Which are the most predictive features according to the training data?***

Using the same metrics for predictiveness as above, the independent variables in this model are **ranked by predictiveness** as follows:

1. foreclosure\_pc\_2010 (t-value: 5.573, log odds: 0.2029)
2. tract\_medage\_2020 (t-value: 3.894, log odds: 0.1455)
3. nhwhite\_2020 (t-value: -4.3319, log odds: -8.6907)
4. mortgaged\_2010 (t-value: 3.677, log odds: 5.1645)
5. avg\_bed (t-value: -2.4329, log odds: -0.6755)
6. pct\_1\_bd (t-value: -2.1151, log odds: -7.0085)
7. medincome\_change\_2010\_2020 (t-value: 1.1874, log odds: 0.6656)
8. pct\_built\_pre\_1960 (t-value: -0.8886, log odds: -1.0883)
9. pct\_built\_2000\_2009 (t-value: 1.3628, log odds: 1.7932)
10. mortgage\_change\_2010\_2015 (t-value: 0.8687, log odds: 2.7022)
11. pop\_change\_pct (t-value: -2.9357, log odds: -2.8016)
12. poverty\_2010 (t-value: -1.926, log odds: -6.7241)

In this case, foreclosure\_pc\_2010 remains the most predictive variable, and tract\_medage\_2020 – which was among the most predictive during the linear regression section – is the second-most predictive variable.

### ***Use the trained model to predict on your testing dataset.***

We used the following code to predict foreclosure\_quantile values based on the values in a test set:


```
# Predict values using the model and test data
predicted_values <- predict(model_6, newdata = test_data, type = "class")

# Calculate the correlation between predicted and actual values
correlation <- cor(as.numeric(predicted_values),
test_data$foreclosure_quantile)
```

For this model, the resulting correlation coefficient between our predicted results and actual values is **0.87**.

This model predicts the exact quantile correctly roughly 60.7% of the time.

Why this model is more accurate than our initial model is unclear.



## B. Naive Bayes:

***Divide your dataset into training and testing set and train the classifier. Report the confusion matrix.***

After dividing our dataset into training and testing sets, we used the following code to train the classifier:

```
# Fit the Naive Bayes model
naive_bayes_model <- naiveBayes(foreclosure_quantile ~ nhwhite_2020 +
                                foreclosure_pc_2010 +
                                tract_medage_2020 +
                                poverty_2010 +
                                medincome_change_2010_2020 +
                                avg_bed +
                                mortgaged_2010 +
                                pct_1_bd +
                                pct_built_pre_1960 +
                                pct_built_2000_2009 +
                                mortgage_change_2010_2015 +
                                pop_change_pct, data =
pg_foreclosures_per_tract_log_reg)
```

And we used the following code to test the classifier:

```
# Predict using the Naive Bayes model
predictions <- predict(naive_bayes_model, newdata = test_data_2)

# Create the confusion matrix
confusion_matrix <- confusionMatrix(predictions,
factor(test_data_2$foreclosure_quantile, levels = 1:5))
```

That produces the following confusion matrix:

Confusion Matrix and Statistics

	Reference				
Prediction	1	2	3	4	5
1	6	2	1	0	0
2	3	2	0	0	1
3	0	1	4	1	2
4	0	0	1	0	3



```
# Fit the Naive Bayes model with Laplace smoothing
naive_bayes_model_laplace <- naiveBayes(foreclosure_quantile ~ nhwhite_2020
+
+ foreclosure_pc_2010 +
+ tract_medage_2020 +
+ poverty_2010 +
+ medincome_change_2010_2020 +
+ avg_bed +
+ mortgaged_2010 +
+ pct_1_bd +
```

```

                                pct_built_pre_1960 +
                                pct_built_2000_2009 +
                                mortgage_change_2010_2015 +
                                pop_change_pct,
                                data =
pg_foreclosures_per_tract_log_reg,
                                laplace = 1)

```

And we used the following code to test the classifier:

```

# Predict using the Naive Bayes model
predictions <- predict(naive_bayes_model_laplace, newdata = test_data_2)

# Create the confusion matrix
confusion_matrix <- confusionMatrix(predictions,
factor(test_data_2$foreclosure_quantile, levels = 1:5))

```

That produces the following confusion matrix:

#### Confusion Matrix and Statistics

	Reference				
Prediction	1	2	3	4	5
1	6	2	1	0	0
2	3	2	0	0	1
3	0	1	4	1	2
4	0	0	1	0	3
5	0	1	1	2	4

#### Overall Statistics

```

Accuracy : 0.4571
95% CI : (0.2883, 0.6335)
No Information Rate : 0.2857
P-Value [Acc > NIR] : 0.02302

```

```

Kappa : 0.3073

```

```

Mcnemar's Test P-Value : NA

```

#### Statistics by Class:

	Class: 1	Class: 2	Class: 3	Class: 4	Class: 5
Sensitivity	0.6667	0.33333	0.5714	0.00000	0.4000

Specificity	0.8846	0.86207	0.8571	0.87500	0.8400
Pos Pred Value	0.6667	0.33333	0.5000	0.00000	0.5000
Neg Pred Value	0.8846	0.86207	0.8889	0.90323	0.7778
Prevalence	0.2571	0.17143	0.2000	0.08571	0.2857
Detection Rate	0.1714	0.05714	0.1143	0.00000	0.1143
Detection Prevalence	0.2571	0.17143	0.2286	0.11429	0.2286
Balanced Accuracy	0.7756	0.59770	0.7143	0.43750	0.6200

Given that all of the categories (quantiles) appear in our training and test data sets, it is unsurprising that the Laplace estimator, which is meant to ensure that categories that do not appear in the training data are still assigned a non-zero probability, does not have an impact on the accuracy of our Naive Bayes model.

### Question 3: Decision Trees and Random Forests

To conduct this form of analysis, we opted to divide our column of foreclosure values per capita in 2020 by three foreclosure risk levels—low, medium (moderate), and high. Our range in values prior to dividing the data was as follows: Min = 8.121113, 25th% = 49.55049, 75th% = 94.04721, and Max = 149.7634. In the screenshot below, we organized this range to fit our various risk levels.

```
pg_foreclosures_per_tract <-
pg_foreclosures_per_tract |>
  mutate(ranges=
    if_else(
      foreclosure_pc_2020 >= 0 & foreclosure_pc_2020 < 49.55049,
      "low",
      if_else(
        foreclosure_pc_2020 >= 49.55049 & foreclosure_pc_2020 < 94.04721,
        "medium",
        "high"
      )
    )
  )
```

From there, we saved a new version of our original dataset called “pg\_fc\_pt” (Prince George’s Foreclosures Per Tract), which removed the “foreclosures per capita in 2020” column to better test the viability of our other variables in the dataset. After completing that step, we set a seed that saved the randomized version of our data. We called this randomized dataset “random\_fc2020” and pulled the 173 rows from our cleaned dataset.

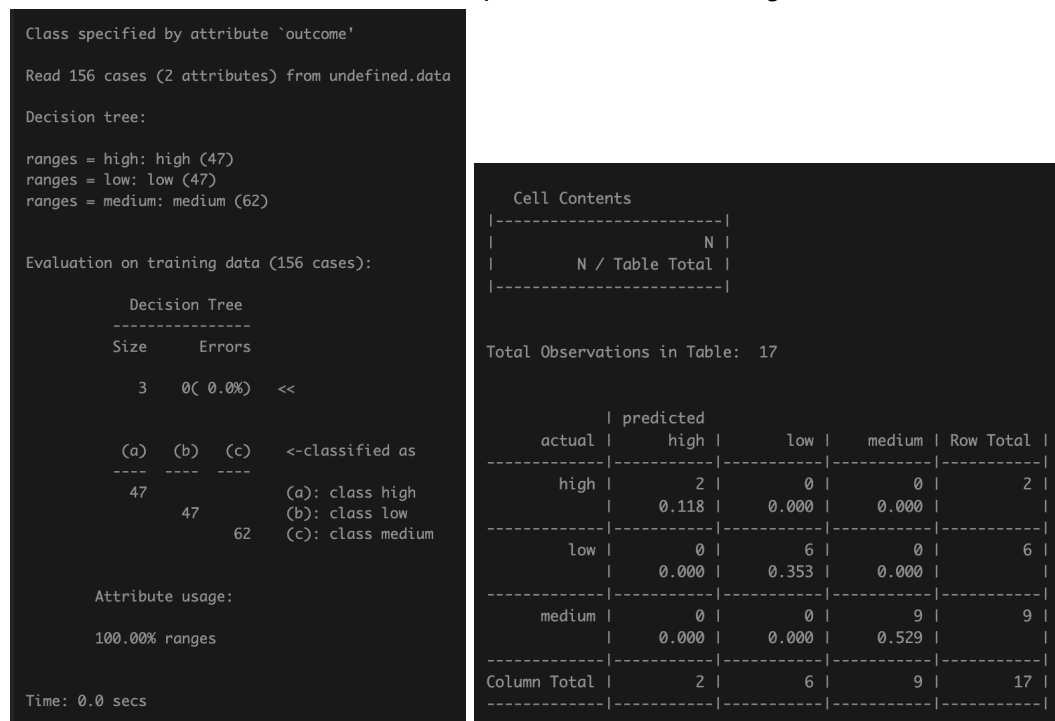
#### A. Testing and Training Datasets

We divided the randomized dataset into testing and training datasets at a 90/10 split of rows. The training dataset contained rows 1 through 156, while the testing dataset contained the remaining rows. These datasets were then entered into proportion tables that enabled us to view their outcome variable distributions based on our previously established risk levels. The

training dataset returned the following distribution: high = 30%, low = 30%, and medium = ~40%. The testing dataset, in comparison, offered a slightly different distribution: high = 11%, low = 35%, and medium = ~53%.

## B. Decision Tree Training and Confusion Matrix

The screenshots below show the results from our initial training on the decision tree and corresponding confusion matrix. The actual values equal the predicted values. This suggests a 100% accuracy rate, but the low sample size also heavily influences this outcome. This also shows that the distribution after the split is similar to the original.



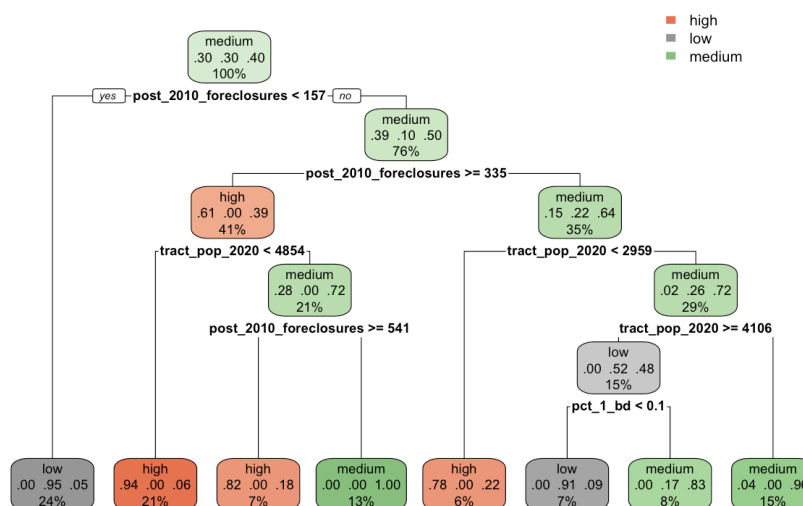
The subsequent screenshot is a plotted example of our decision tree, visualizing how the model determines outcomes using our low, medium, and high structure. We initially expected variables like median income to be our ideal predictors, as opposed to the variables we ultimately observed in the tree. We were surprised to see that omission considering the variable's relevance in other forms of predictive analysis. We didn't necessarily expect post-2010 foreclosures as a predictive variable. However, this makes some sense in retrospect because it is likely that tracts that previously showed lower numbers of foreclosures might be more likely to have a lower foreclosure rate in the future compared to other tracts with a higher number of foreclosures. Tract populations in 2020 were another interesting variable to us. The tree indicates that higher population counts in a given tract led to higher foreclosure potential. This may be explained by the fact that a highly populated area presents more opportunities for foreclosures as opposed to a less populated area. The other unexpected variable was the percentage of 1 bedroom units that were less than 10%. This resulted in a lower probability of foreclosing, which could have been the case because residents of one-bedroom properties

might have relatively lower monthly payments and living expenses compared to families or residents of properties with a higher number of bedrooms. However, we would need to examine this further with more data.

Our tree begins by looking at post-2010 foreclosure values per tract in our dataset that are less than 157. If "yes," the tree predicted a low chance of foreclosure, with 24% of the dataset's values. If "no," it checked post-2010 foreclosures greater than or equal to 335. "Yes" values to that variable predicted a high chance of future foreclosures, and it branched off to check for 2020 tract populations of less than 4,854. A "yes" to that variable indicates a high chance of foreclosure, and 21% of our values followed this path. A "no" to 2020 tract populations less than 4,854 led to examining post-2010 foreclosures greater than or equal to 541. "Yes" suggested a high chance of foreclosure, with 7% of our values falling in this range, while "no" predicted a medium chance of foreclosure with 13% of our values.

Going back to post-2010 foreclosures greater than or equal to 335, a "no" examined 2020 tract populations of less than 2,959. A "yes" predicted a high chance of foreclosure and showed 6% of our values in this outcome, while a "no" led to looking at 2020 tract populations greater than or equal to 4,106. "Yes" outcomes led to a low chance of foreclosure (15%) and split off again based on the percentage of 1-bedroom units that were less than 10%. For this split, we saw 7% of our "yes" values with a low chance to foreclose, while "no" values (8%) indicated a medium chance to foreclose. Finally, "no" answers to 2020 tract populations greater than or equal to 4,106 (15%) indicated a medium chance to foreclose.

In summary, we observed a relatively even distribution of outcomes in the decision tree. Roughly 31% fell under low foreclosure probability outcomes, while 36% and 34% fell under medium and high foreclosure probability outcomes, respectively. The most significant low probability indicator came from tracts with post-2010 foreclosure values less than 157 (24%), the most significant medium outcomes came from 2020 tract populations greater than or equal to 4,106 (15%), and the most significant high foreclosure probability outcome came from 2020 tract populations less than 4,854 (21%).



### C. Boosting

The results shown in the screenshot below indicate that boosting was attempted with three trials. However, the process was abandoned after the first trial because of the high level of accuracy achieved due to the low number of classifiers. This highlighted a limitation of our data, given that more data is needed in order to fully complete the boosting process and draw valuable conclusions from this step.

```
Call:
C5.0.default(x = training[46], y = training$ranges, trials = 3)

C5.0 [Release 2.07 GPL Edition]      Sat Apr  6 08:36:34 2024
-----

Class specified by attribute 'outcome'

Read 156 cases (2 attributes) from undefined.data

----- Trial 0: -----

Decision tree:

ranges = high: high (47)
ranges = low: low (47)
ranges = medium: medium (62)

*** boosting reduced to 1 trial since last classifier is very accurate
*** boosting abandoned (too few classifiers)

Evaluation on training data (156 cases):

      Decision Tree
      -----
      Size      Errors
      3      0( 0.0%) <<

      (a) (b) (c) <-classified as
      --- --- ---
      47      47      62
              (a): class high
              (b): class low
              (c): class medium

Attribute usage:

100.00% ranges
```

### D. Bagging and Random Forests

Like in previous steps, we set the seed and randomized the dataset. Then, we conducted the bagging analysis, which added bootstrap replications. We ran 50 bootstrap replications as it is a high number that should yield adequate results. After running the code, the misclassification error was 22.5%, which indicates our accuracy was 77.5%. This suggests that random forests with bagging help raise our level of accuracy. However, this analysis would also greatly benefit from more data to run through these models.

Bagging classification trees with 50 bootstrap replications

```
Call: bagging.data.frame(formula = ranges ~ ., data = bagfc2020, nbagg = 50,
  coob = TRUE, control = rpart.control(minsplit = 2, cp = 0,
  min_depth = 2))
```

Out-of-bag estimate of misclassification error: 0.2254

As it pertains to our random forest, we first trained the model with four different numbers of trees. Then, we set a prediction outcome after running the test data through the random forest. Our range of accuracy shown in the screenshot below suggests that the model has moderate predictive performance. However, this finding does not necessarily suggest the format is sufficient for predictive analysis.

In this section, we also calculated importance scores for our dataset variables. Post-2010 foreclosures, mortgaged 2020, owner-occupied properties in 2015, owner-occupied properties in 2020, and the percentage of properties built pre-1960 in each tract had relatively higher importance across all trees in the random forest.

Overall, we would most likely want a different model to help us predict foreclosures per capita.

```
# training
random_forest = randomForest(as.factor(train_data$ranges)~.,data = train_data,ntree = 4,importance = T)

# predictions about the test data
predictions <- predict(random_forest, test_data)

importance_scores <- importance(random_forest)
# we used this function to ascertain the model's accuracy.
accuracy <- sum(predictions == test_data$ranges) / nrow(test_data)
print(paste("Accuracy:", round(accuracy, 2)))
print(importance_scores)
```

```
[1] "Accuracy: 0.69"
```

	high	low	medium	MeanDecreaseAccuracy	MeanDecreaseGini
avg_bed	0.00000000	0.000000	0.0000000	0.00000000	0.0000000
tract_number	0.00000000	0.000000	1.9985407	1.99973582	0.9630376
tract_homevalue_2020	-0.06730536	0.000000	-1.1547005	-0.45883147	3.0162659
tract_medage_2020	1.15470054	1.756917	1.1547005	1.56414491	3.3428104
tract_medincome_2020	0.00000000	0.000000	0.0000000	0.00000000	0.0000000
tract_medincome_2010	0.00000000	0.000000	0.0000000	0.00000000	0.6469697
tract_pop_2010	0.00000000	1.154701	-1.1547005	1.15470054	2.7215097
tract_pop_2020	1.15470054	0.000000	0.0000000	1.15470054	1.9495042
great_recession_foreclosures	1.15470054	1.946657	1.1547005	1.67977499	2.9313724
post_2010_foreclosures	5.37471543	1.644573	0.7821117	3.23700172	17.9559473

#### **Question 4: Comparative Analysis. (Jasper)**

[5 pts] Write a summary of all classifiers, their predictive quality and which one would you use to answer your research question(s).

Please add a Contributions section at the end of your report clearly stating the contributions of each student to the report, coding and presentation preparation and recording (Youtube video). For example, in a team with three members VFM, LM and AM, you could write "Question 1a: developed by VFM and LM; Question 1b: developed by VFM; Question 1c: developed by VFM and AM; Question 1d: developed by All; Question 1e: developed by VFM. VFM prepared 70% of the code; LM worked on 20% of the code and AM on 10%. All students contributed equally to the preparation and recording of the presentation".

Team efforts only work if all team members distribute the workload equitably. If a team member's contributions are much lower than the others' I will prorate the final grade accordingly.

#### **Q4: Comparative analysis**

Team 2's overarching research question:

Is it possible to forecast the per capita rate of foreclosure for a census tract in Prince George's County, Maryland, using identified independent variables?

Foreclosure is a cross cutting issue that bears negative implications within the financial system writ large, but also affects populations at the community and family level. According to a study conducted by the Urban Institute in Washington DC, the process and experience of foreclosure has generational ripple effects that may stem from legal exposure during the foreclosure process, displacement and housing instability, prolonged financial insecurity, personal/familial stress and mental hardship.

For these reasons, Team 2 has approached it's comparative analysis through two lenses:

1. Which classifier is most equipped to answer our primary research question from the standpoint of statistical accuracy
2. Which classifier bears the ability to supplement decision making capabilities of a decision maker in the public or private sector concerned with mitigating the effects of foreclosure in an area of interest?

Due to the extensive ecological, economical, and demographic impacts of foreclosure, Team 2 believes public and private sector decision makers have vested interest in working to identify areas within their respective locales, states, or countries that are most at risk for foreclosure as to formulate mitigative strategies geared to cater communities classified at high risk. Team 2, toward answering our identified research question, constructed several classifiers that could be used to predict per capita foreclosure rates as well as predict tracts most at risk for foreclosure.



## Classifiers and Result Summaries:

### Linear Regressions

Team 2's linear regressions model leveraged the following variables determined to be most predictive amongst our available after regularization data:

1. pct\_built\_2000\_2009 with a coefficient of 64.77035.
2. pct\_built\_2010\_2019 with a coefficient of 55.26635.
3. avg\_bed with a coefficient of 51.21755.
4. pct\_built\_1990\_1999 with a coefficient of 29.19928.
5. pct\_built\_1980\_1989 with a coefficient of 12.14104.
6. ownoccupied\_2015 with a coefficient of 78.21088.
7. pct\_1\_bd with a coefficient of -103.1094.
8. pct\_built\_2020\_later with a coefficient of 1.904033.
9. mortgage\_change\_2015\_2020 with a coefficient of 32.56724.
10. ownoccupied\_change\_2010\_2015 with a coefficient of -47.27385.

### Multivariate Linear regressions

The most predictive features of Team 2's Multivariate Linear regressions were medincome\_change\_2010\_2015, medincome\_change\_2015\_2020, poverty\_2010, medincome\_change\_2010\_2020, and nhwhite\_2020. This model had a correlation value of 79% between predicted and real values, and had a mean squared error of 436.931.

### Logistic Regression and Naïve Bayes

#### Logistic Regression:

Team 2's logistic regression model leveraged 13 variables that were determined to be the most predictive amongst our existing variable set. These variables are noted in detail within the Logistic Regression section above in this Milestone 2 report.

Of these 13 variables, nhwhite\_2020, the percentage of residents in a tract claiming non-Hispanic White ancestry, as well as foreclosure\_pc\_2010 were most predictive. After executing this model against our test dataset, we found that the model accurately predicted per capita quantile foreclosure rates at the tract level 60% of the time, measured against quantiles of our dependent variable foreclosure\_pc\_2020.

#### Naïve Bayes:

Team 2 trained its Naïve Bayes model using the same set of selected variables used to train its logistic regression model described in this report. Team 2's Naïve Bayes classifier was accurate 45% of the time according to the confusion matrix produced by the model, though it was significantly more accurate than the no-information prediction.

## Decision Trees

Team 2, in constructing its decision tree model, established low, medium and high ranges for prioritized variables, which were post-2010 foreclosures, 2020 tract population and percentage of 1 bedroom units. This breakdown essentially allowed us to classify census tracts within Prince Georges County MD, based on foreclosure risk level (low, medium, or high) based on our selected attributes within our decision tree.

Team 2's decision tree model ultimately reported 31% of Prince George's County census tracts fell under low foreclosure risk, 36% fell under medium risk, and 34% fell under high risk, based on our identified variables of interest. This decision tree model reported its results with 100% accuracy; however, the small size of the tract-level data sample used by Team 2 likely influenced this accuracy level. Team 2 would expect potentially richer insights from block-level data, which would allow decision makers to administer mitigative efforts to target areas with higher accuracy.

## Random Forest:

Team 2's Random Forest model was 77.5% accurate, with a misclassification error of 22.5% after executing bagging models with 50 bootstrap replications. This model was trained using four separate sets of trees for training.

## Outcome

Team 2's most accurate classifiers were its decision tree and its logistic regression classifiers. Based on these results we would likely want to use a logistic regression model for the specific purpose of predicting per capita foreclosures based on our most predictive variables. Such a classifier has the potential to answer our research question of whether we can predict per capita foreclosures at the census tract level in Prince George's County. Adding additional, granular data from the block level of our Prince George's county census tracts could potentially improve the accuracy and utility of our overall model. Team 2 aims to explore this in future milestones of our project.

To inform a decision maker, we would likely opt to use a combination of methods. Team 2 believes the decision tree classifier it has constructed could be used effectively to classify census tracts based on their foreclosure risk. In conjunction, the logistic regression model could be used to forecast foreclosure rates. Such a combination of data streams would effectively enable a decision maker to identify areas of risk/concern, with a forecast to what unmitigated outcomes could look like. Based on Team 2's research question and working data, this combination of products would create a machine learning pipeline that would be the most useful to a decision maker.

## Potential Drawbacks

Decision trees have a high degree of interpretability as they visually represent decision making and are relatively intuitive to understand and visualize. For decision makers and those

responsible for briefing and or influencing decision makers, these attributes are extremely important. Generally, when it comes to the accuracy of decision trees, the models are prone to overfitting, which typically occurs when there are too few data points. A decision tree with a model with higher utility may need to incorporate more data points to overcome this. Logistic regressions additionally require sufficient sample sizes for reliable results to occur. If this condition is not met, logistic regression models additionally are likely to overfit predictions.

**Contributions:**

**Question 1:** Krehl Kasayan and Paul Kiefer

**Question 2:** Paul Kiefer

**Question 3:** Pablo Suarez

**Question 4:** Jasper Evans

**Report:** Equal input by all four members.

**Video Presentation:** Video editing by Pablo Suarez